

Technical Task: Social Feed Explorer

Time budget: 48 hours from receipt

Context

Our product is a social platform with feeds that render large volumes of content, with search and filtering across that content. This task is a small, self-contained version of that problem, so we can see how you approach the kinds of challenges you'd actually face working with us.

What you're building

A feed browser app that works with a large dataset (5,000–10,000 items) of posts and/or users. You'll fetch, display, search, and filter this data in a responsive UI.

Dataset / mock API: Use any mock API/Data. You are not expected to build a backend.

Requirements

1. Data & list rendering

- Load and display the dataset as a scrollable list/feed.
- The list must remain smooth and responsive at full scale (thousands of items) — think carefully about how you render this.
- Implement pagination, infinite scroll, or “load more” (your choice).

2. Search

- Add a search input that filters the list by relevant fields (e.g. name/content).
- It should feel responsive and stay usable at full data scale. How you achieve that is up to you.

3. Filters

- Add at least 2–3 filter dimensions (e.g. date range, category/tag, status) that can be combined with each other and with search.
- Filter/search state should be sharable.

4. State management

- Use whatever state management approach(es) you think are appropriate — you're free to combine local state, context, a state library, and a data-fetching library however you see fit.
- We're interested in *why* you split state the way you do, not which specific tools you reach for.

5. Styling

- Build with Tailwind. Responsive layout. Doesn't need to be pixel-perfect, but should show reasonable design taste.

Optional (stretch — only if you have time left)

- Optimistic UI update on some action (e.g. a “like” button with rollback on failure)
- Skeleton/loading states
- Error handling/retry logic
- Keyboard navigation/accessibility basics
- Light and dark theme
- A short `NOTES.md` on what you'd do differently at 10x the data scale, or what you deliberately skipped and why

What to submit

- A link to a repo (GitHub or similar)
- A README with:
 - Setup/run instructions
 - Any assumptions you made
 - What you'd improve with more time
 - Anything you intentionally left out and why

If anything in the brief is ambiguous, make a reasonable assumption, note it in your README, and move on — we want to see how you handle ambiguity, not a flood of clarifying questions.

Good luck, and thank you for the time you put into this.